

Hyperparameter Learning, Linear Regression with Nonnegative Target, Logistic Regression

Krisztian Buza

Department of Artificial Intelligence
Eötvös Loránd University
Budapest, Hungary

Announcements

Announcements

- Tentative Grading Scheme

85 pts – 100 pts	→ 5 („very good“)
70 pts – 84 pts	→ 4 („good“)
55 pts – 69 pts	→ 3 („average“)
40 pts – 54 pts	→ 2 („sufficient“)
0 pts – 39 pts	→ 1 („not sufficient“)

Announcements


- Quick Test Replacement Test (QTRT)
 - In the same session with the „big test“
 - 11th Dec: „big test“ (60 minutes) + ~~QTRT (30 minutes)~~
 - Due to multiple requests, QTRT will be organized between 16th Dec and 20th Dec
 - You may collect maximal 14 points which replace the result of the 7 worst quick tests

Students' Presentations

- Everyone in the team should contribute to the presentation
- Team members can distribute the work internally as they wish (e.g., it is not needed that everyone talks, but everyone should do something for the presentation)
- First slide: title of the presentation
name and Neptun code of all team members
- Last slide: who did what
- Send your slides in PDF(!) format to buza@inf.elte.hu at least two days before the presentation

Overfitting from the Perspective of RMSE

Calculation of Root Mean Square Error




y	\hat{y}
549	-770
100	-470
976	1132
441	1433
3401	3435
2600	4537
4241	5638
7150	6539
9100	7541
9000	8542

$n = 10$


$$\frac{1}{10} \left((549 - (-770))^2 + (100 - (-470))^2 + \right. \\ \left. + (976 - 1132)^2 + (441 - 1433)^2 + \right. \\ \left. + (3401 - 3435)^2 + (2600 - 4537)^2 + \right. \\ \left. + (4241 - 5638)^2 + (7150 - 6539)^2 + \right. \\ \left. + (9100 - 7541)^2 + (9000 - 8542)^2 \right)$$

Calculation of Root Mean Square Error




y	\hat{y}
549	254
100	293
976	805
441	958
3401	2433
2600	3602
4241	5004
7150	6332
9100	7999
9000	9867

Calculation of Root Mean Square Error




y	\hat{y}
549	534
100	131
976	608
441	894
3401	3100
2600	2951
4241	4011
7150	7247
9100	9081
9000	9002


Root Mean Square Error of the Models



1085.88




745.50




248.66

Root Mean Square Error on New Data




y	\hat{y}
9	-1171
3464	4337
9125	8041
10609	9043

Root Mean Square Error on New Data






y	\hat{y}
9	4779
3464	2979
9125	7204
10609	37347

Root Mean Square Error on New Data



y	\hat{y}
9	4779
3464	2979
9125	7204
10609	37347

Root Mean Square Error of the Models



on training data:

1085.88

745.50

248.66

on new data (test data):

1202.28

209.73

13616.00

Hyperparameter Learning

Parameters and Hyperparameters

- Polynomial regression: $\hat{y} = w^{(0)} + w^{(1)}x + w^{(2)}x^2 + \dots + w^{(p)}x^p$

Parameters: $w^{(0)}, w^{(1)}, \dots, w^{(p)}$

- Hyperparameters: parameters of the learning algorithm (optimization algorithm), such as
 - Degree of the polynomial
 - Learning rate
 - „Importance“ of the Regularization Term...
- Parameters are determined by the learning algorithm, whereas hyperparameter may be set by the „expert“

How to Find the Appropriate Values of Hyperparameters?

Hyperparameter Learning

- Determining the best (appropriate) values of hyperparameters is part of the training process.
- Therefore, we need **completely new** data („Test data 2“) in order to have an unbiased („fair“) estimate of the quality (RMSE...) of the model with „best“ hyperparameters.
- Terminology:

Training data	→	Training data
Test data 1	→	Validation data
Test data 2	→	Test data

(sometimes „validation data“ refers to „test data 2“ and „test data“ is used to refer to „test data 1“)

Batch Gradient Descent and Stochastic Gradient Descent

Gradient Descent for Linear Regression

(Regularization is omitted for simplicity)

- We consider models of the form $\hat{y} = \sum_{j=1}^m w^{(j)} x^{(j)}$
- Procedure:
 - **Step 1** Set $w^{(1)}, w^{(2)}, \dots, w^{(m)}$ to some random values
 - **Step 2** for j in $1 \dots m$

*The gradient of the
objective function
(a.k.a. cost function)*

$$w^{(j)} \leftarrow w^{(j)} - \epsilon \frac{1}{n} \sum_{i=1}^n 2x_i^{(j)} (\hat{y}_i - y_i)$$

where $\hat{y}_i = \sum_{j=1}^m w^{(j)} x_i^{(j)}$

- **Step 3** Repeat Step 2 as long as you can non-negligibly decrease the error

Variants of Gradient Descent

- Gradient is calculated using all instances of the training data → **gradient descent**
- Gradient is calculated using a subset of the instances of the training data → **batch gradient descent**
 - Batch size: the number of instances in that subset that is used to calculate the gradient
 - Each time when step 2 is executed, a different subset should be selected
- Gradient is calculated using a single instance → **stochastic gradient descent**

Stochastic Gradient Descent for Lin. Reg.

(Regularization is omitted for simplicity)

- We consider models of the form $\hat{y} = \sum_{j=1}^m w^{(j)} x^{(j)}$
- Procedure:
 - **Step 1** Set $w^{(1)}, w^{(2)}, \dots, w^{(m)}$ to some random values
 - **Step 2** $x_i \leftarrow$ **select a random instance of the training data**
for j in $1 \dots m$

$$w^{(j)} \leftarrow w^{(j)} - \epsilon \frac{1}{n} \sum_{i=1}^n 2x_i^{(j)} (\hat{y}_i - y_i)$$

$$\text{where } \hat{y}_i = \sum_{j=1}^m w^{(j)} x_i^{(j)}$$

- **Step 3** Repeat Step 2 as long as you can non-negligibly decrease the error

Batch Gradient Descent for Lin. Reg.

(Regularization is omitted for simplicity)

- We consider models of the form $\hat{y} = \sum_{j=1}^m w^{(j)} x^{(j)}$
- Procedure:
 - **Step 1** Set $w^{(1)}, w^{(2)}, \dots, w^{(m)}$ to some random values
 - **Step 2** $D' \leftarrow$ select a random subset of the training data for j in $1 \dots m$


$$w^{(j)} \leftarrow w^{(j)} - \epsilon \frac{1}{n} \sum_{x_i \text{ in } D'} 2x_i^{(j)} (\hat{y}_i - y_i)$$

$$\text{where } \hat{y}_i = \sum_{j=1}^m w^{(j)} x_i^{(j)}$$

- **Step 3** Repeat Step 2 as long as you can non-negligibly decrease the error

How do Different Versions of Gradient Descent Find the Optimum?

Blue: gradient descent
Red: stochastic gradient descent



Linear Regression with Nonnegative Target


Linear Regression with Nonnegative Target

- We consider models of the form

$$z = \sum_{i=0}^m w^{(i)} x^{(i)}$$

$$\hat{y} = \begin{cases} z & \text{if } z \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

Linear Regression with Nonnegative Target



Cost Function and Its Partial Derivatives

- Cost function: mean of squared errors (regularisation is omitted for simplicity)

$$E = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

\uparrow

$$\hat{y} = \begin{cases} z & \text{if } z \geq 0 \\ 0 & \text{otherwise} \end{cases} \quad z = \sum_{i=0}^m w^{(i)} x^{(i)}$$

- Partial Derivative w.r.t. $w^{(j)}$

$$\frac{\partial E}{\partial w^{(j)}} = \begin{cases} \frac{1}{n} \sum_{i=1}^n 2x_i^{(j)} (\hat{y}_i - y_i) & \text{if } z \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

Cost Function and Its Partial Derivatives

- Cost function: mean of squared errors (regularisation is omitted for simplicity)

$$E = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

$$\hat{y} = \begin{cases} z & \text{if } z \geq 0 \\ 0 & \text{otherwise} \end{cases} \quad z = \sum_{i=0}^m w^{(i)} x^{(i)}$$

- Partial Derivative w.r.t. $w^{(j)}$ **calculated on a single instance**

$$\frac{\partial E}{\partial w^{(j)}} = \begin{cases} \frac{1}{n} \sum_{i=1}^n 2x_i^{(j)} (\hat{y}_i - y_i) & \text{if } z \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

Stochastic Gradient Descent for Linear Regression with Nonnegative Target (Regularization is omitted for simplicity)

- Procedure:
 - **Step 1** Set $w^{(1)}, w^{(2)}, \dots, w^{(m)}$ to some random values
 - **Step 2** $x_i \leftarrow$ **select a random instance of the training data**
for j in $1 \dots m$
$$w^{(j)} \leftarrow w^{(j)} - \epsilon \frac{\partial E}{\partial w^{(j)}}$$
 - **Step 3** Repeat Step 2 as long as you can non-negligibly decrease the error

Linear Regression with Nonnegative Target

- The gradient may be zero for some instances with substantial error.
- Therefore, the previous algorithm may find a model that fits some of the instances well whereas some other instances are ignored.
- What about using several models „together“?

Regression vs. Classification


Regression vs. Classification

- **Regression**: target is continuous
 - price of an apartment, product, etc.
 - UPDRS score
- **Classification**: target is discrete (it takes the value from a finite set)
 - Recognition of handwritten digits
 - Is the e-mail spam or not?
(class 0: „spam“, class 1: „not spam“)
- Regressor: the model used to solve regression tasks
- Classifier: the model used to solve classification tasks

Logistic Regression

(it is a classifier, not a regressor!)

„Classification via Regression“



„Classification via Regression“



Hypothesis

$$\mathbf{x} = (x^{(0)}, x^{(1)}, \dots, x^{(m)})$$

$$x^{(0)} = 1$$

$$x^{(i)} \in \mathbf{R}$$


$$\mathbf{w} = (w^{(0)}, w^{(1)}, \dots, w^{(m)})$$

$$w^{(i)} \in \mathbf{R}$$

- Linear Regression: $\hat{y} = \mathbf{w}\mathbf{x}$
- Logistic Regression: $\hat{y} = \sigma(\mathbf{w}\mathbf{x})$

Sigmoid function
(a.k.a. logistic function)

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$



- Probabilistic interpretation

Objective Function

- Linear Regression (L_2 regularization):

$$E = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2 + \lambda \sum_{i=1}^p (w^{(i)})^2$$

- Can we use the „same“ function with $\hat{y} = \sigma(\mathbf{w}\mathbf{x})$?


Objective Function – Logistic Regression

- Cost for the i -th instance:

$$\text{cost}(\hat{y}^{(i)}, y^{(i)}) = \begin{cases} -\log(\hat{y}^{(i)}) & \text{if } y^{(i)} = 1 \\ -\log(1 - \hat{y}^{(i)}) & \text{if } y^{(i)} = 0 \end{cases}$$

- Total cost
(for a set of instances):

$$\frac{1}{n} \sum_{i=1}^n \text{cost}(\hat{y}^{(i)}, y^{(i)})$$



Objective Function – Logistic Regression

- Cost for the i -th instance:

$$\begin{aligned} \text{cost}(\hat{y}^{(i)}, y^{(i)}) &= \begin{cases} -\log(\hat{y}^{(i)}) & \text{if } y^{(i)} = 1 \\ -\log(1 - \hat{y}^{(i)}) & \text{if } y^{(i)} = 0 \end{cases} \\ &= -y^{(i)}\log(\hat{y}^{(i)}) - (1 - y^{(i)})\log(1 - \hat{y}^{(i)}) \end{aligned}$$

- Total cost
(for a set of instances):

$$-\frac{1}{n} \sum_{i=1}^n (y^{(i)}\log(\hat{y}^{(i)}) + (1 - y^{(i)})\log(1 - \hat{y}^{(i)}))$$

- This cost function is also known as *cross-entropy*.

Cross-entropy vs. Mean Squared Error

See **Figure 5** in

X Glorot, Y Bengio (2010):
Understanding the difficulty of training deep feedforward neural
networks

http://proceedings.mlr.press/v9/glorot10a/glorot10a.pdf?hc_location=ufi

Summary

Summary

- Selection of appropriate hyperparameters is part of the training process
- Batch gradient descent, stochastic gradient descent
- Linear Regression with Nonnegative Target
- Classification
- Logistic Regression

Essential Concepts

- Parameter – Hyperparameter
- Stochastic Gradient Descent – Batch Gradient Descent
- Batch Size
- Training data – Validation Data – Test data
- Regression – Classification
- Regressor – Classifier
- Class
- Cross-entropy