

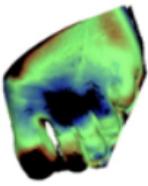
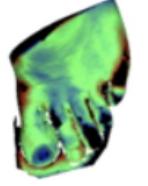
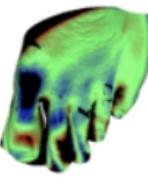
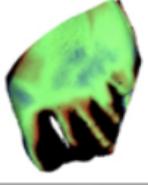
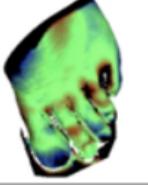
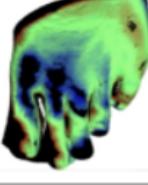
Deep Learning in a Nutshell, Convolutional Neural Networks, Implementation of Machine Learning Algorithms

Krisztian Buza

Department of Artificial Intelligence
Eötvös Loránd University
Budapest, Hungary

Announcement

Possible project / thesis topic: who made which statue?

	Sosikles Type (Capitolino)	Sciarra Type (Copenhagen)	Mattei Type (Tivoli)
Doryphoros (Pompeii Statue)			
Diadoumenos (Athens)			
Doryphoros (Pompeii Statue)			
Diadoumenos (Athens)			



Pheidias ? Polykleitos ? Krésilas ?



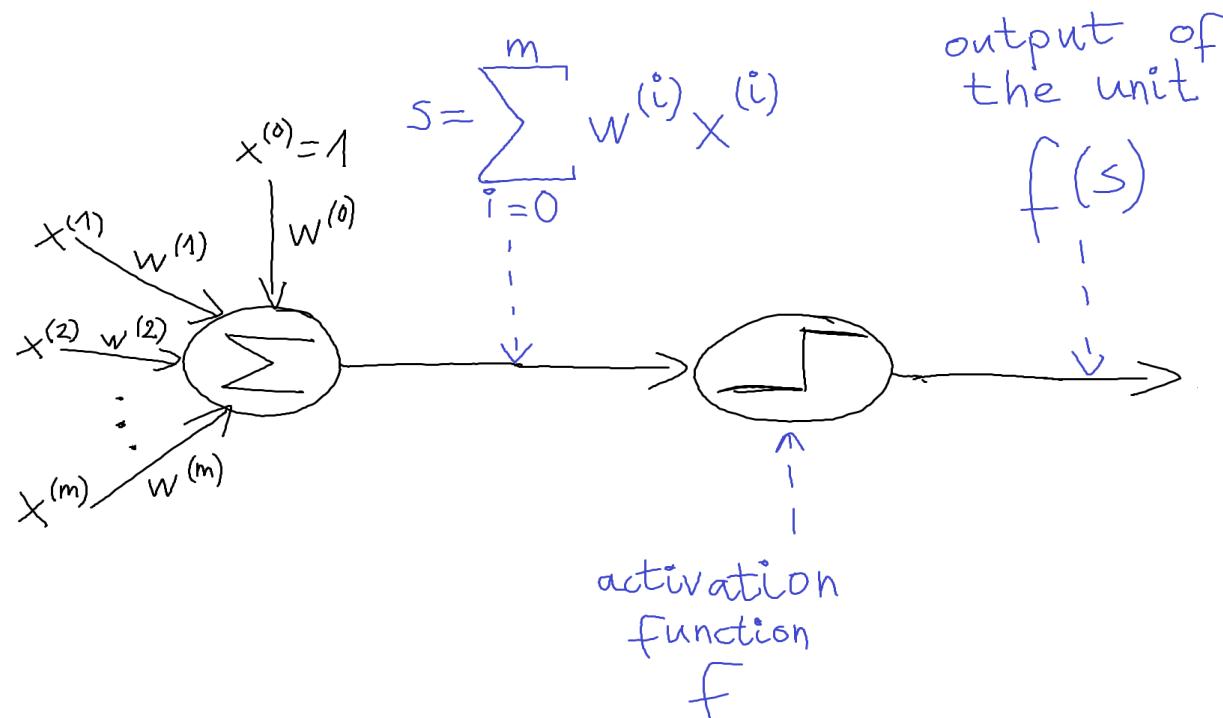
DORYPHOROS



DIADOU MENOS

Deep Learning in a Nutshell

A Neural Unit



$x^{(1)}, x^{(2)}, \dots, x^{(m)}$ = inputs of the unit

$w^{(1)}, w^{(2)}, \dots, w^{(m)}$ = weights of $x^{(1)}, x^{(2)}, \dots, x^{(m)}$

$w^{(0)}$ = bias weight

„What was wrong with backpropagation in 1986?“

(Geoff Hinton, „Deep Learning“, May 22, 2015)

- Our labeled datasets were thousands of times too small.
- Our computers were millions of times too slow.
- We initialized the weights in a stupid way.
- We used the wrong type of non-linearity.

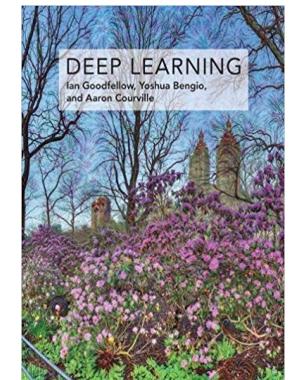
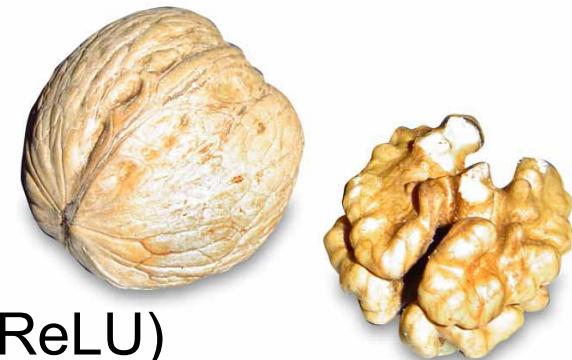


Image: https://upload.wikimedia.org/wikipedia/commons/3/34/Geoffrey_Hinton_at_UBC.jpg
Eviatar Bach [CC BY-SA 3.0 (<https://creativecommons.org/licenses/by-sa/3.0/>)]

Deep Learning in a Nutshell

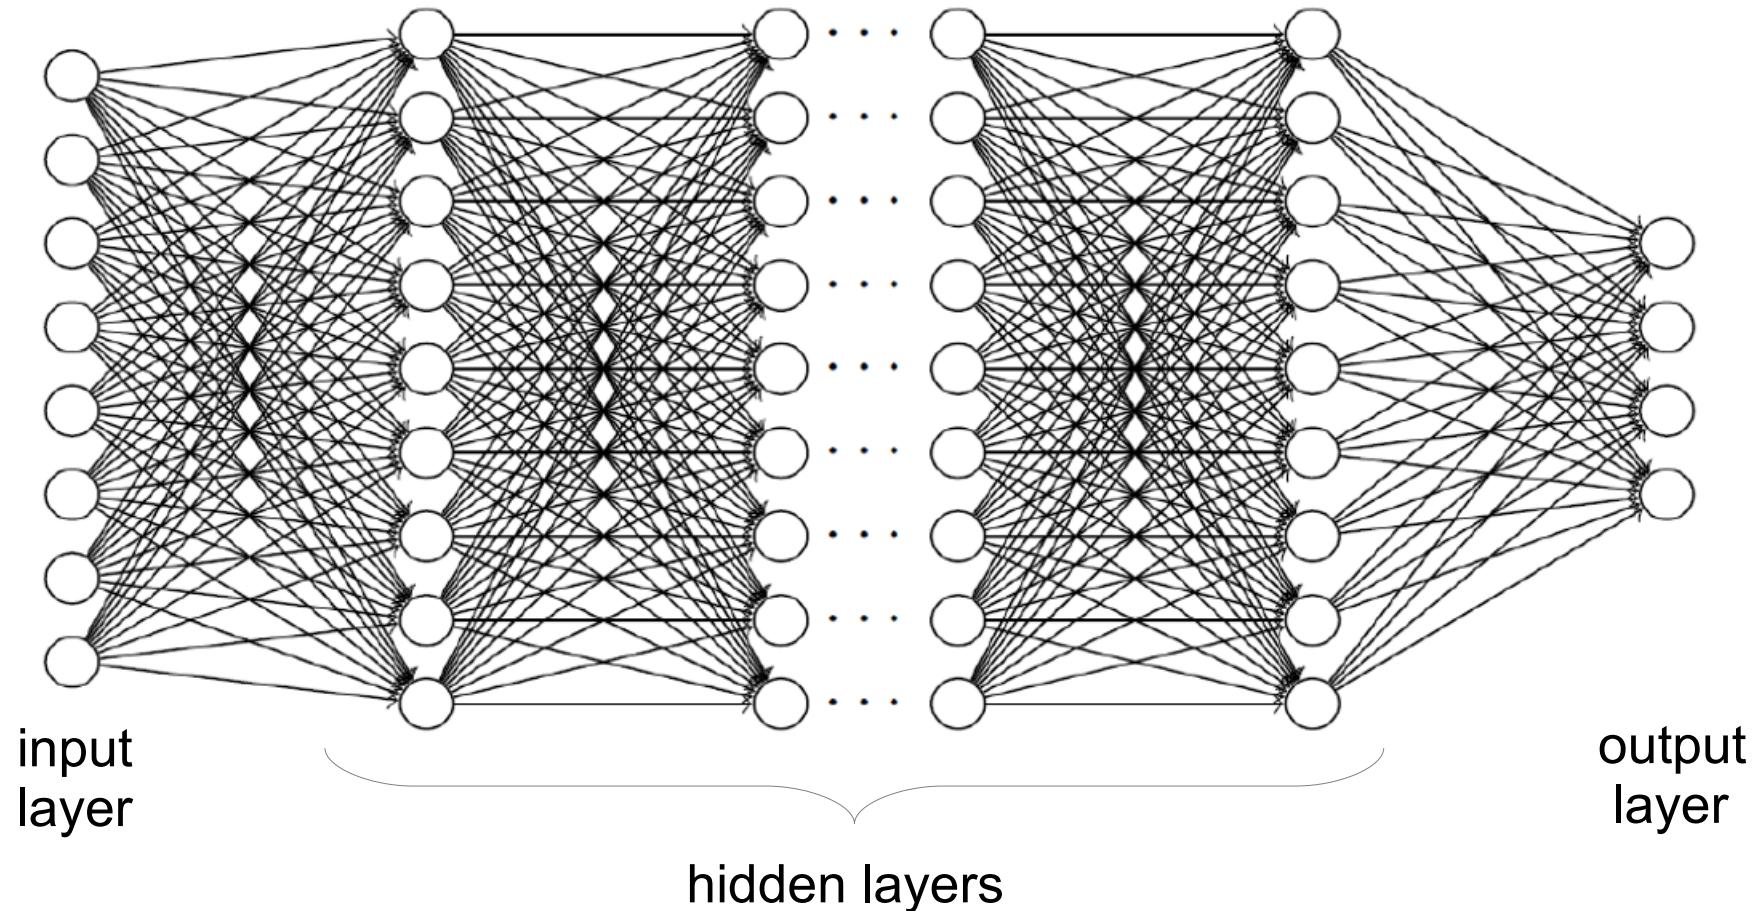
- Size and structure of the network:
few layers → many layers
- Activation function: sigmoid → rectified linear unit (ReLU)
- Loss function: quadratic loss → cross-entropy
- Initialization of weights: random → (unsupervised) pre-training
- Size of training data, much more memory,
distributed computation (HPC), GPUs, TPUs...
- New regularization techniques:
“sparsity-enforcing” regularisation terms,
drop-out, early stop,
parameter sharing (convolutional networks)
- New optimization techniques (e.g. ADAM)
- Data augmentation

By Horst Frank - photo taken by Horst Frank, CC BY-SA 3.0.
<https://commons.wikimedia.org/w/index.php?curid=14209>



<http://www.deeplearningbook.org>

Deep Feed-Forward Neural Networks



Loss Function: Cross-entropy vs. Mean Squared Error

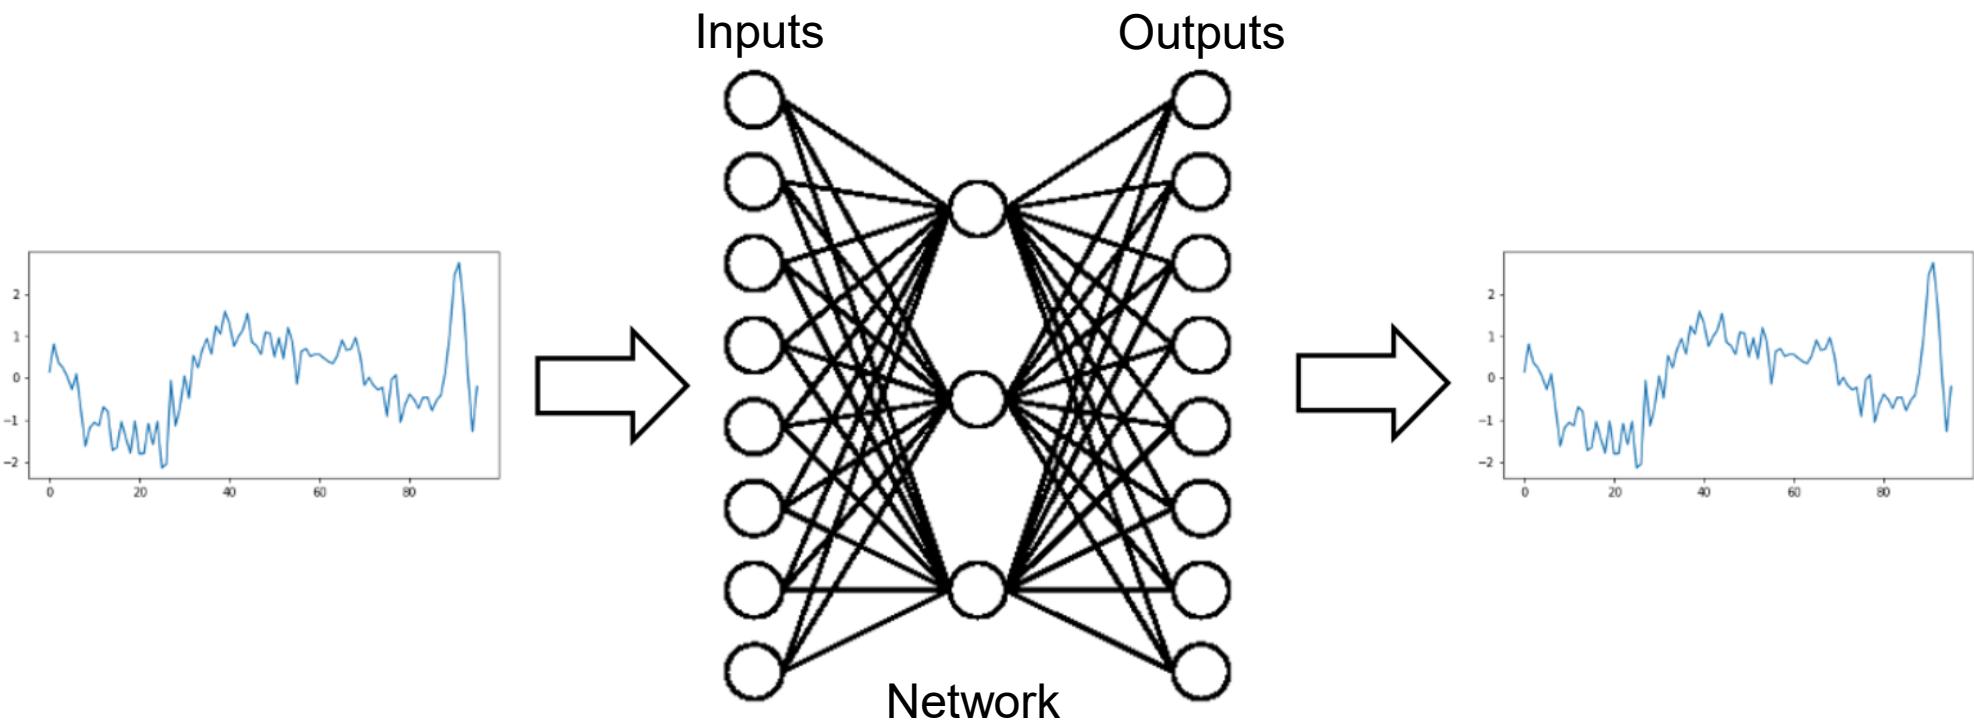
See **Figure 5** in

X Glorot, Y Bengio (2010):
Understanding the difficulty of training deep feedforward neural
networks

<http://proceedings.mlr.press/v9/glorot10a/glorot10a.pdf>

Initialisation of the Weights

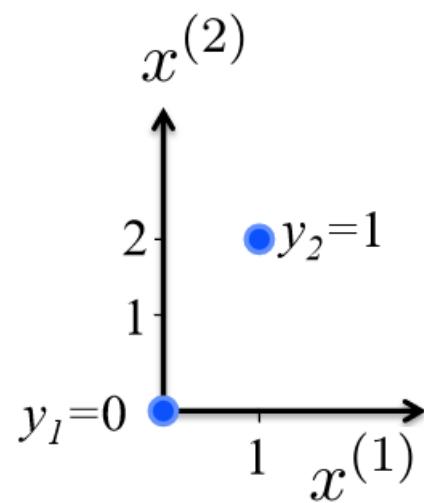
- Unsupervised pre-training: autoencoders
- Supervised pre-training (transfer learning):
 - Train a network for a different (but somehow related...) task
 - Re-use some of the weights
(e.g. weights of the first few convolutional layers)



Sparsity-enforcing (encouraging) Regularisation

- In the example below, all the three models below have the same prediction performance (on training data)

- L_2 regularisation: $\sum_{i=1}^p (w^{(i)})^2$
- L_1 regularisation: $\sum_{i=1}^p |w^{(i)}|$
(sparsity-encouraging)



$$\hat{y} = 1x^{(1)} + 0x^{(2)}$$

$$0^2 + 1^2 = 1$$

$$|0| + |1| = 1$$

$$\hat{y} = 0x^{(1)} + \frac{1}{2}x^{(2)}$$

$$0^2 + \left(\frac{1}{2}\right)^2 = \frac{1}{4}$$

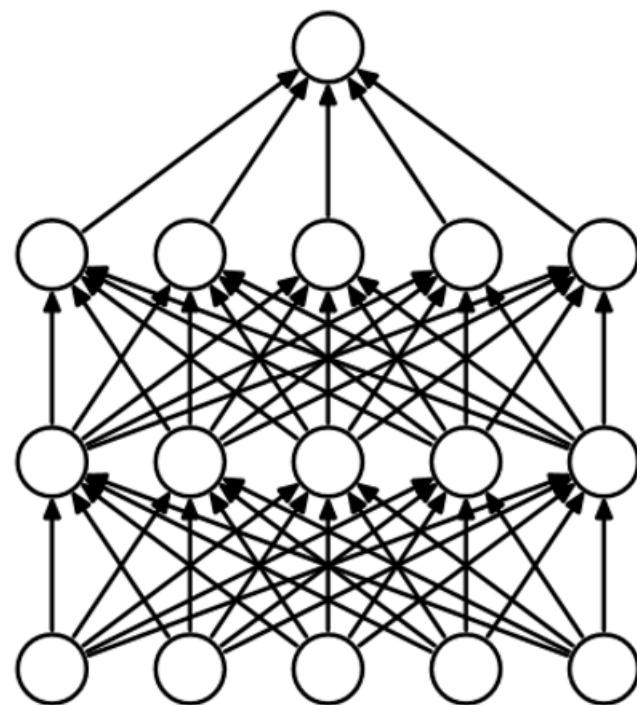
$$|0| + \left|\frac{1}{2}\right| = \frac{1}{2}$$

$$\hat{y} = \frac{1}{3}x^{(1)} + \frac{1}{3}x^{(2)}$$

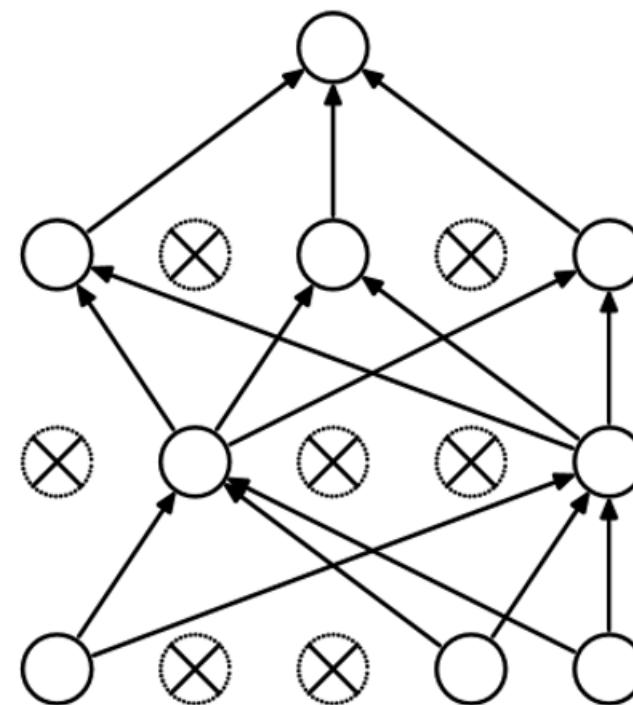
$$\left(\frac{1}{3}\right)^2 + \left(\frac{1}{3}\right)^2 = \frac{2}{9}$$

$$\left|\frac{1}{3}\right| + \left|\frac{1}{3}\right| = \frac{2}{3}$$

Dropout



(a) Standard Neural Net

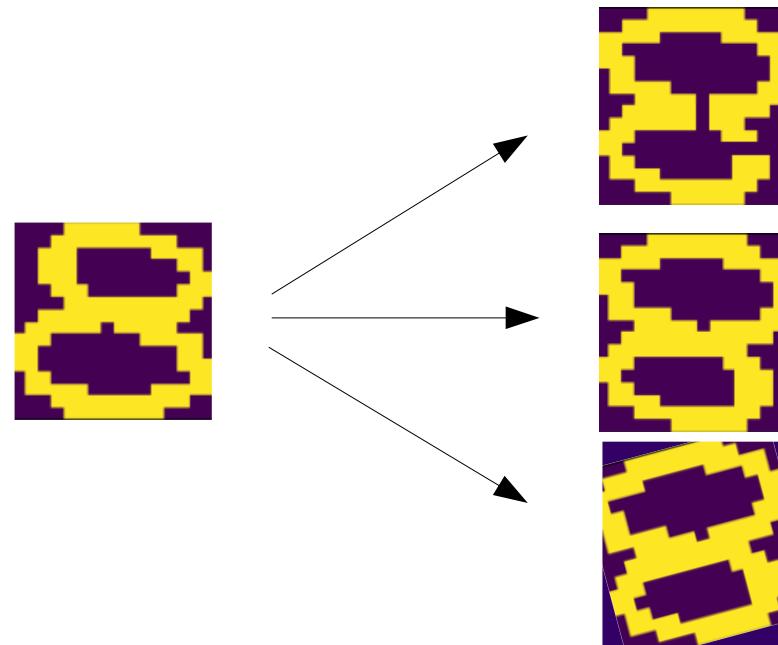


(b) After applying dropout.

Srivastava et al. (2014): Dropout: A Simple Way to Prevent Neural Networks from Overfitting,
Journal of Machine Learning Research

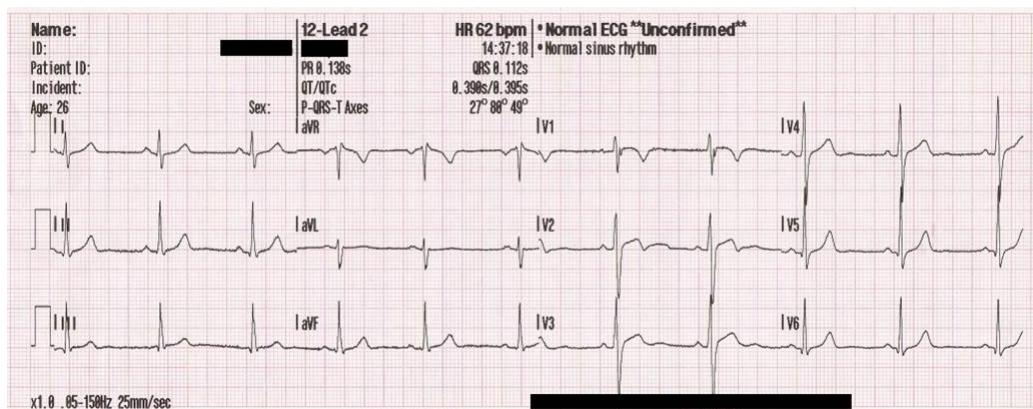
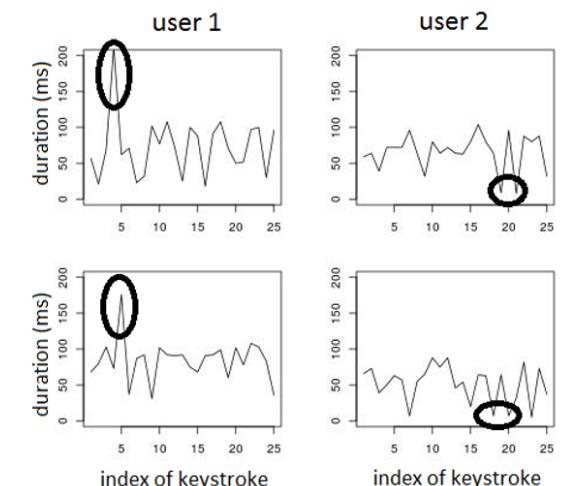
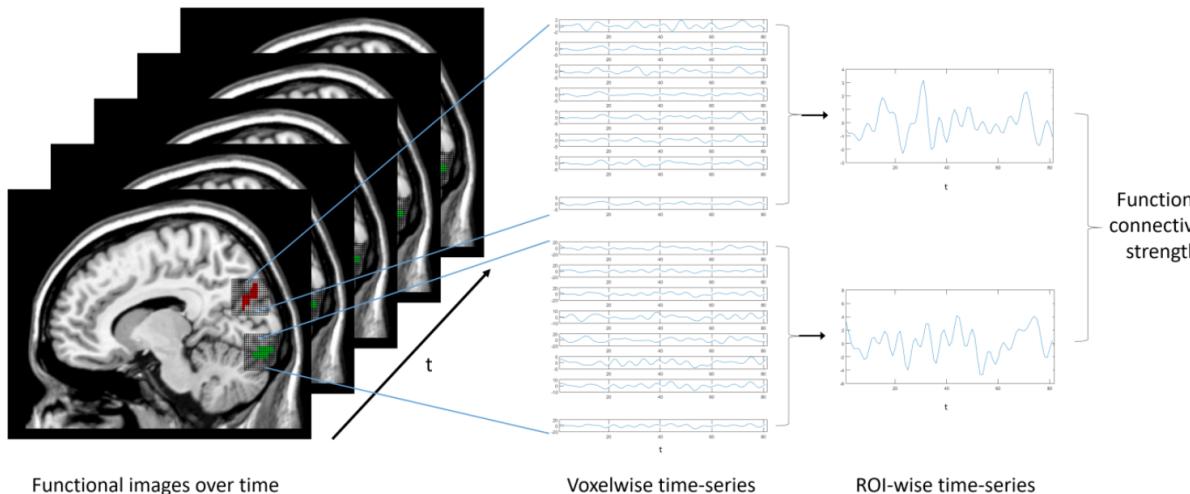
Data Augmentation

- „Generate“ additional training data
(usually from existing training instances)
- Images: add noise, rotate, zoom-in/zoom-out...



Convolutional Networks

Time Series Classification – Examples



Images in the bottom, from left to right:

By MoodyGroove - 2007-01-24 (original upload date) Original uploader was MoodyGroove at en.wikipedia, Public Domain, <https://commons.wikimedia.org/w/index.php?curid=5266589>

By Thuglas at English Wikipedia - Transferred from en.wikipedia to Commons by Sreejithk2000 using CommonsHelper., Public Domain, <https://commons.wikimedia.org/w/index.php?curid=10827060>

By JSquish - Own work, CC BY-SA 3.0, <https://commons.wikimedia.org/w/index.php?curid=16181727>

Convolution*

Input of the convolution (time series):

-0.8	-0.5	-0.2	0.2	0.6	0.8	0.9	1.0	0.9	0.7	0.2	-0.3	-0.9	-0.2	0.5	0.6
------	------	------	-----	-----	-----	-----	-----	-----	-----	-----	------	------	------	-----	-----

* Remark: while being conceptually the same as traditional convolution in mathematics, the convolution used in neural networks is slightly different in terms of technical details.

Convolution*

Input of the convolution (time series):

-0.8	-0.5	-0.2	0.2	0.6	0.8	0.9	1.0	0.9	0.7	0.2	-0.3	-0.9	-0.2	0.5	0.6
------	------	------	-----	-----	-----	-----	-----	-----	-----	-----	------	------	------	-----	-----

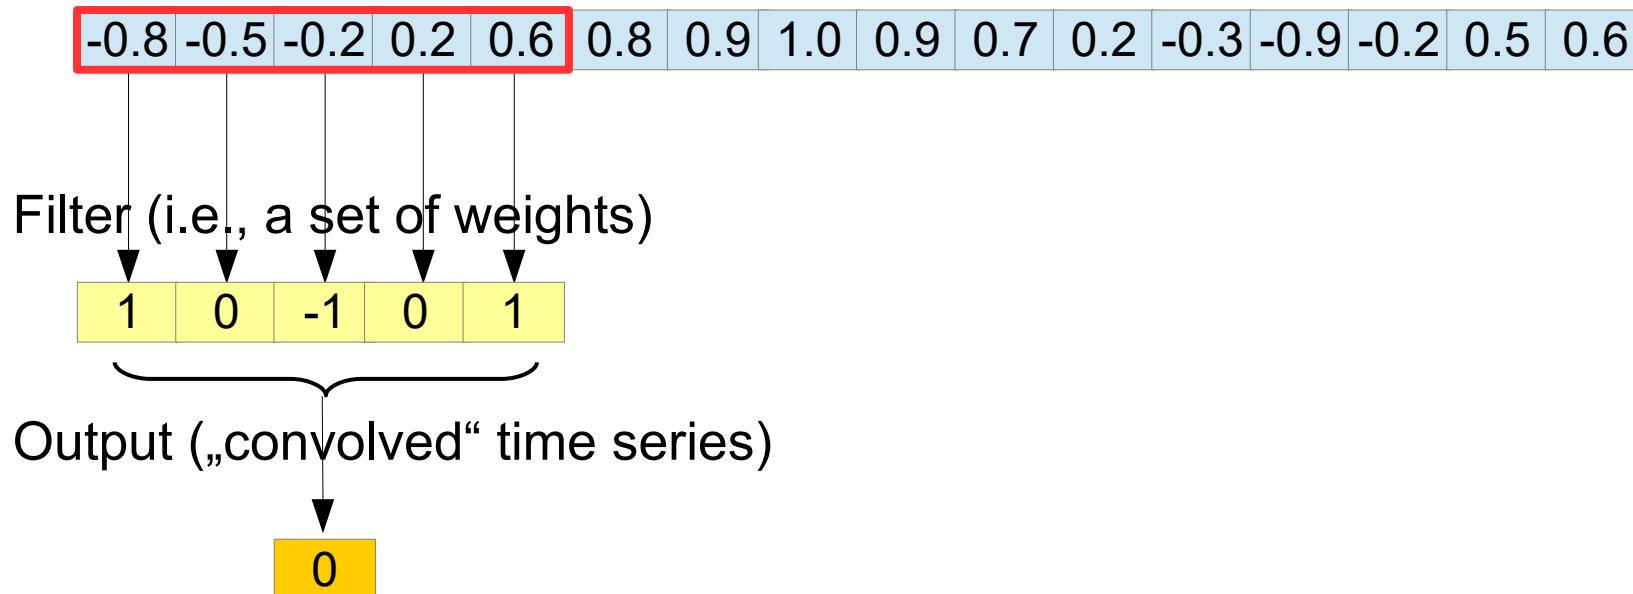
Filter (i.e., a set of weights)

1	0	-1	0	1
---	---	----	---	---

* Remark: while being conceptually the same as traditional convolution in mathematics, the convolution used in neural networks is slightly different in terms of technical details.

Convolution*

Input of the convolution (time series):

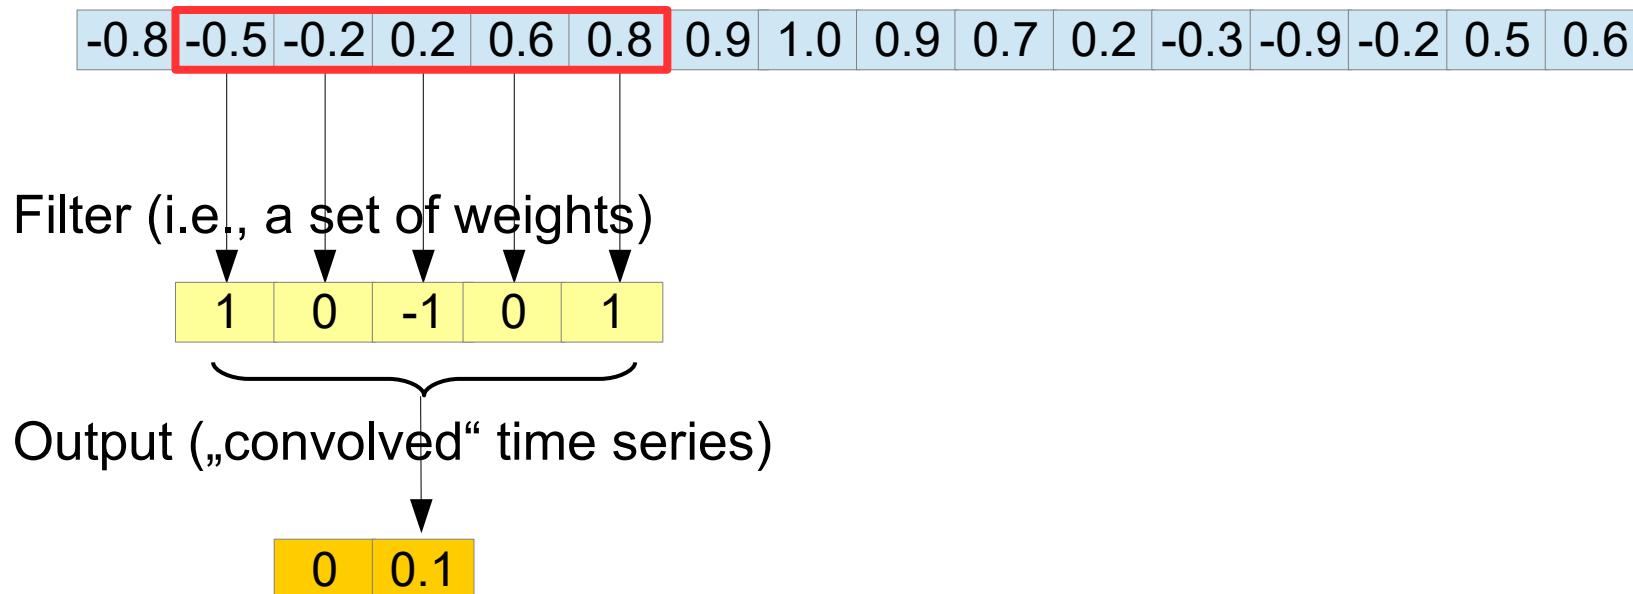


$$(-0.8) \times 1 + (-0.5) \times 0 + (-0.2) \times (-1) + 0.2 \times 0 + 0.6 \times 1 = 0$$

* Remark: while being conceptually the same as traditional convolution in mathematics, the convolution used in neural networks is slightly different in terms of technical details.

Convolution*

Input of the convolution (time series):



$$(-0.5) \times 1 + (-0.2) \times 0 + 0.2 \times (-1) + 0.6 \times 0 + 0.8 \times 1 = -0.1$$

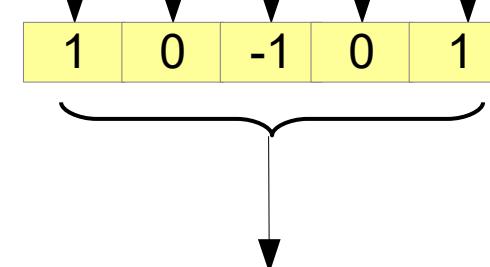
* Remark: while being conceptually the same as traditional convolution in mathematics, the convolution used in neural networks is slightly different in terms of technical details.

Convolution*

Input of the convolution (time series):



Filter (i.e., a set of weights)



Output („convolved“ time series)



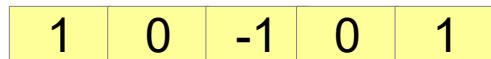
* Remark: while being conceptually the same as traditional convolution in mathematics, the convolution used in neural networks is slightly different in terms of technical details.

Convolution*

Input of the convolution (time series):



Filter (i.e., a set of weights)



Output („convolved“ time series)



* Remark: while being conceptually the same as traditional convolution in mathematics, the convolution used in neural networks is slightly different in terms of technical details.

Convolution*

Input of the convolution (time series):

0	0	-0.8	-0.5	-0.2	0.2	0.6	0.8	0.9	1.0	0.9	0.7	0.2	-0.3	-0.9	-0.2	0.5	0.6	0	0
0	0	0.9	0.3	0.1	-0.2	0.5	0.3	0.1	0	0.2	-0.1	-0.2	0.4	0.5	0.5	0.6	0.3	0	0

Filter (i.e., a set of weights)

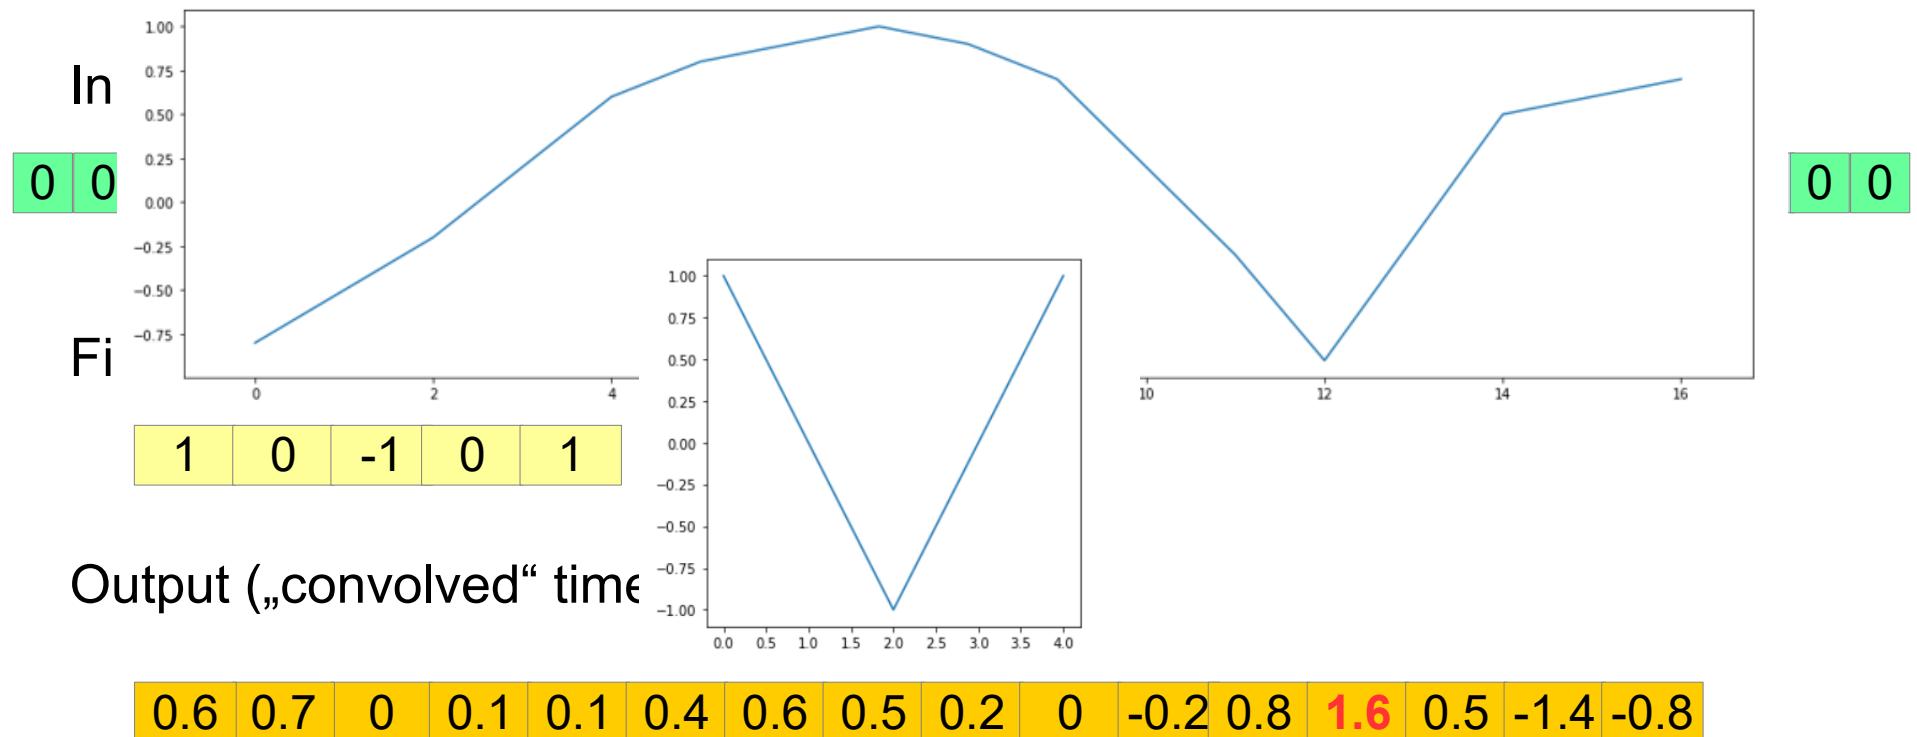
1	0	-1	0	1
0.5	0.3	0.1	-0.2	-0.3

Output („convolved“ time series)

0.6	1.0	0.4	0.1	0.1	0.5	0.9	0.7	0.4	0	-0.4	0.5	1.4	0.7	-1.0	-0.3
-----	-----	-----	-----	-----	-----	-----	-----	-----	---	------	-----	-----	-----	------	------

* Remark: while being conceptually the same as traditional convolution in mathematics, the convolution used in neural networks is slightly different in terms of technical details.

Convolution*



Output („convolved“ time)

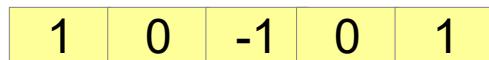
* Remark: while being conceptually the same as traditional convolution in mathematics, the convolution used in neural networks is slightly different in terms of technical details.

Convolution and Max Pooling*

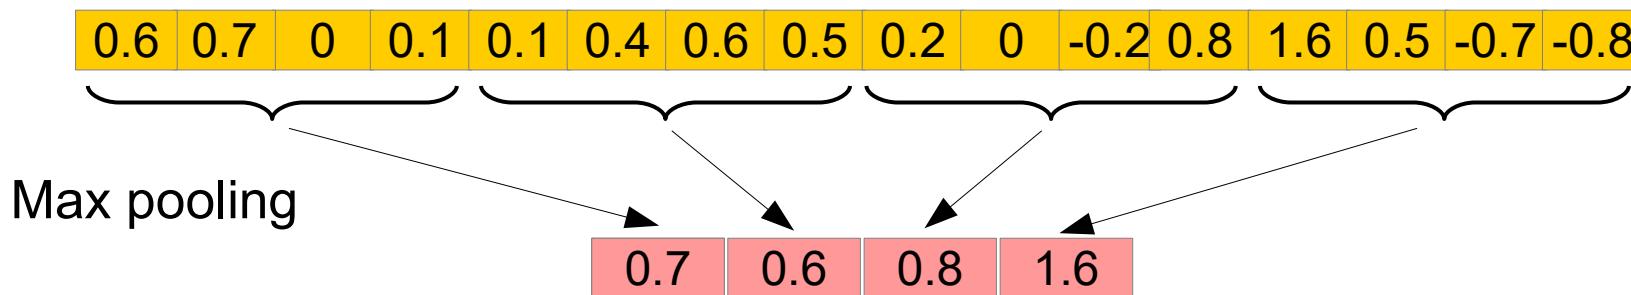
Input of the convolution (time series):



Filter (i.e., a set of weights)

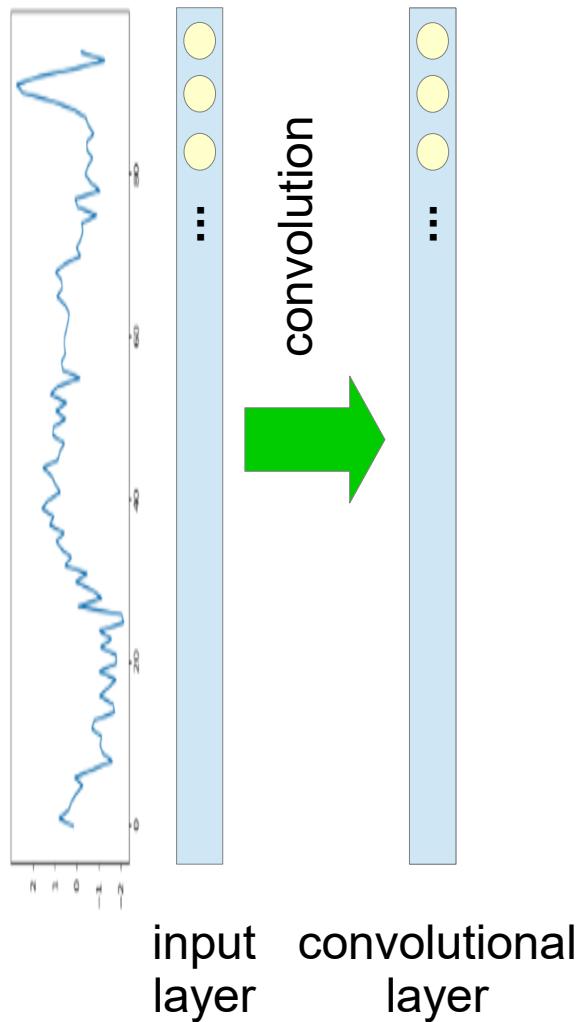


Output („convolved“ time series)

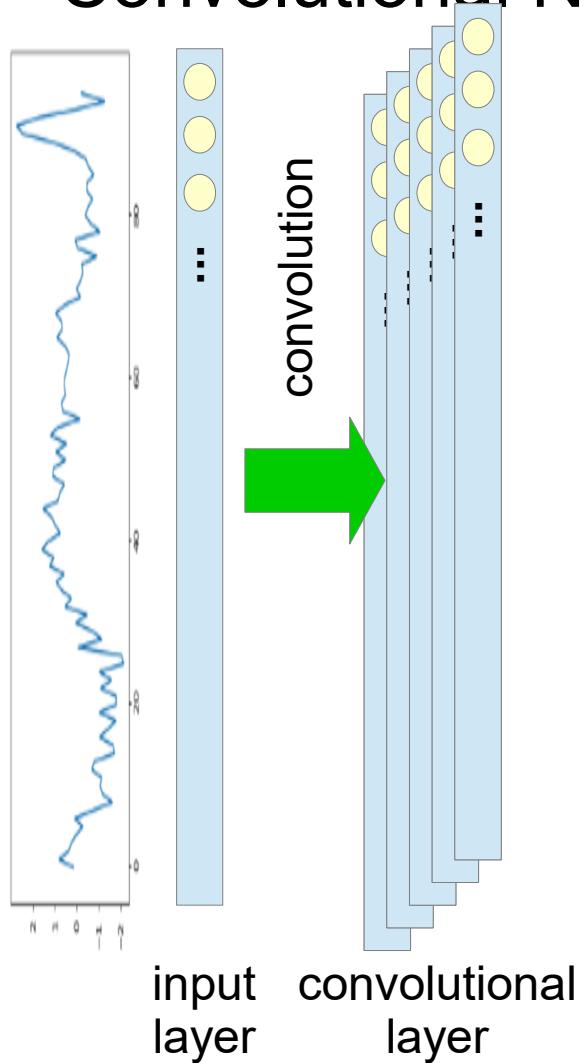


* Strictly speaking, max pooling has nothing to do with convolution, however, in convolutional neural networks (CNNs), the convolutional layer is often followed by a max pooling layer.

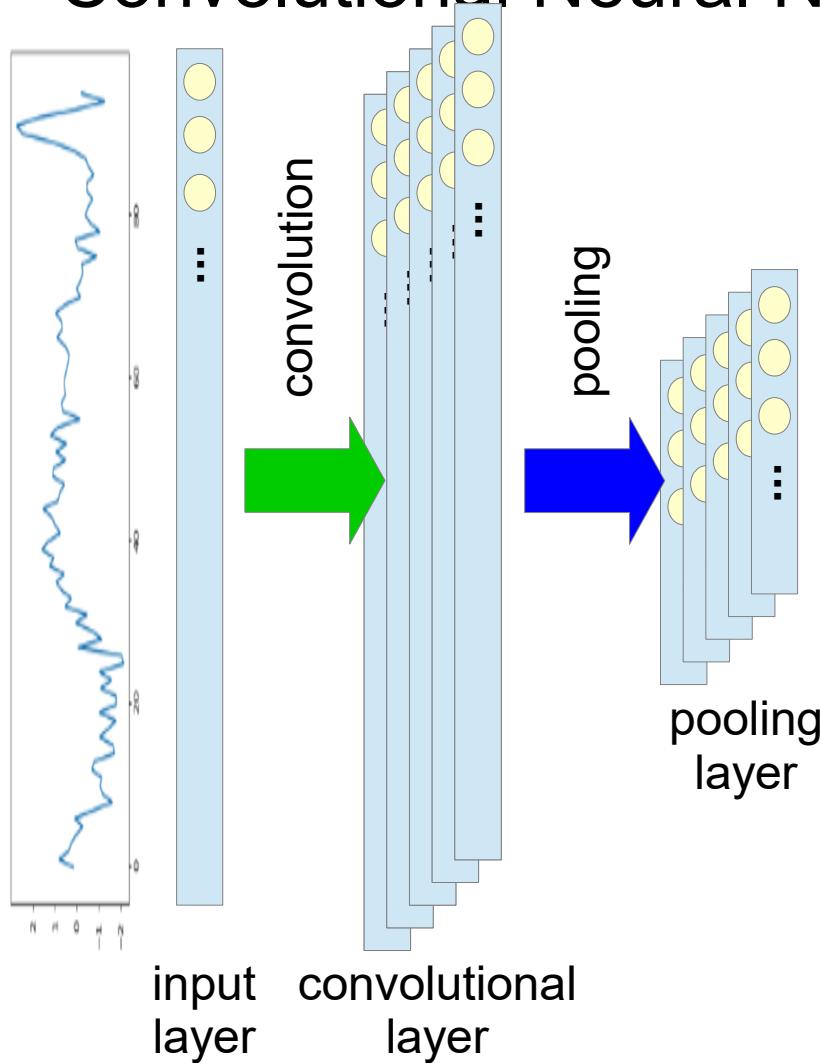
Convolutional Neural Networks



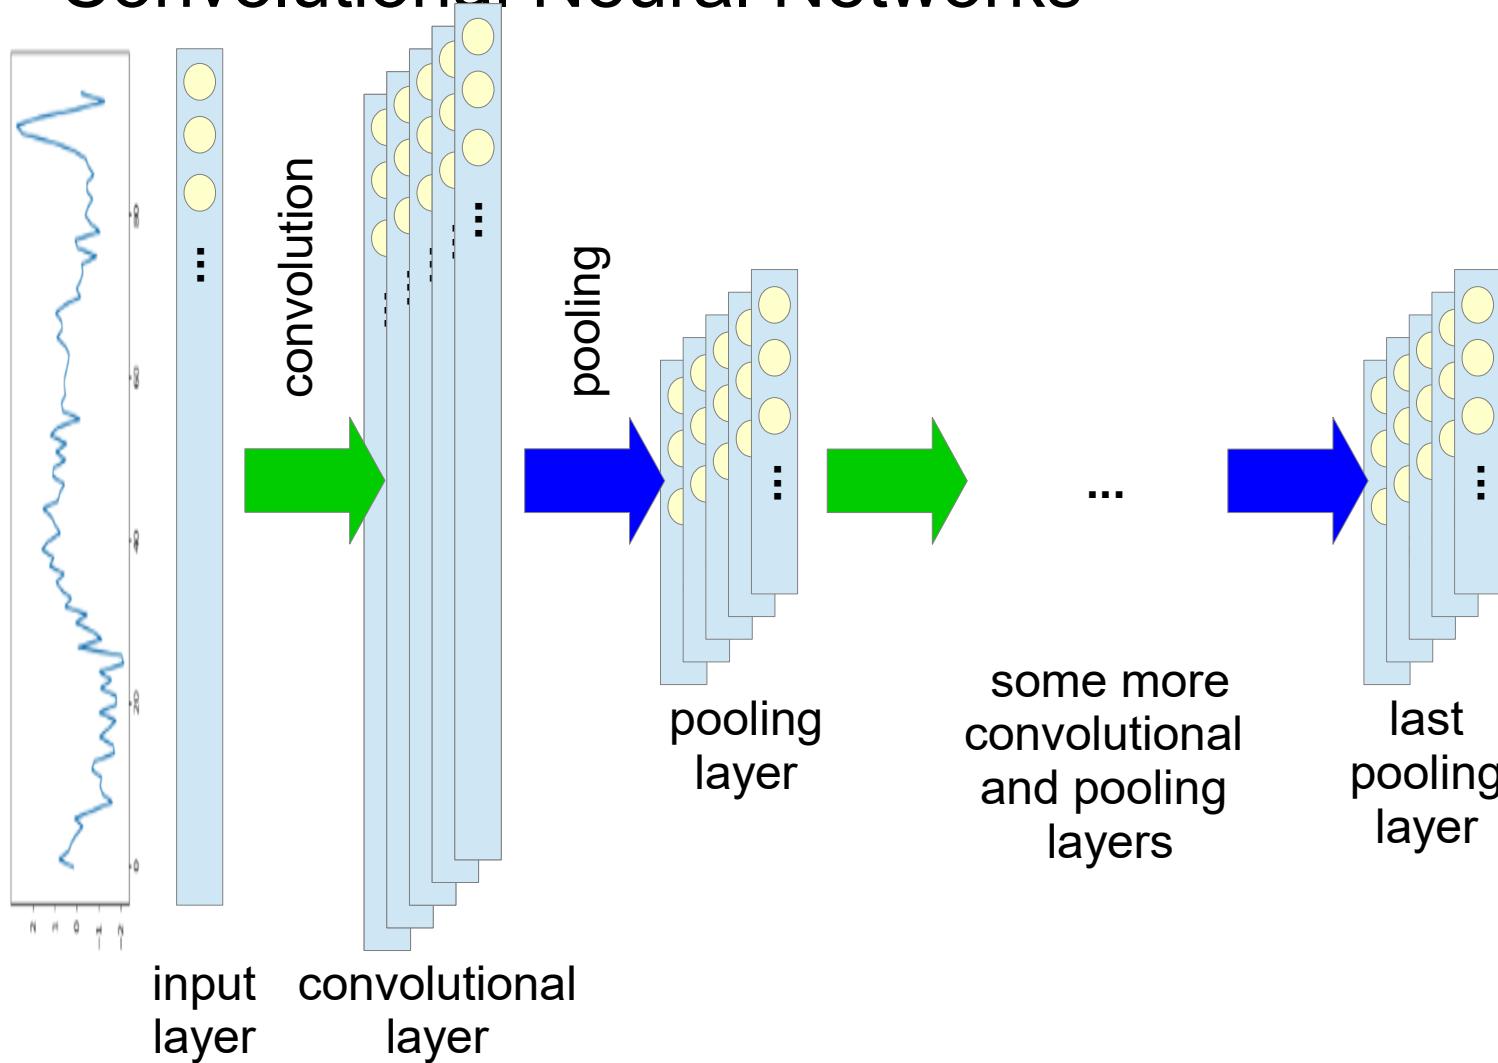
Convolutional Neural Networks



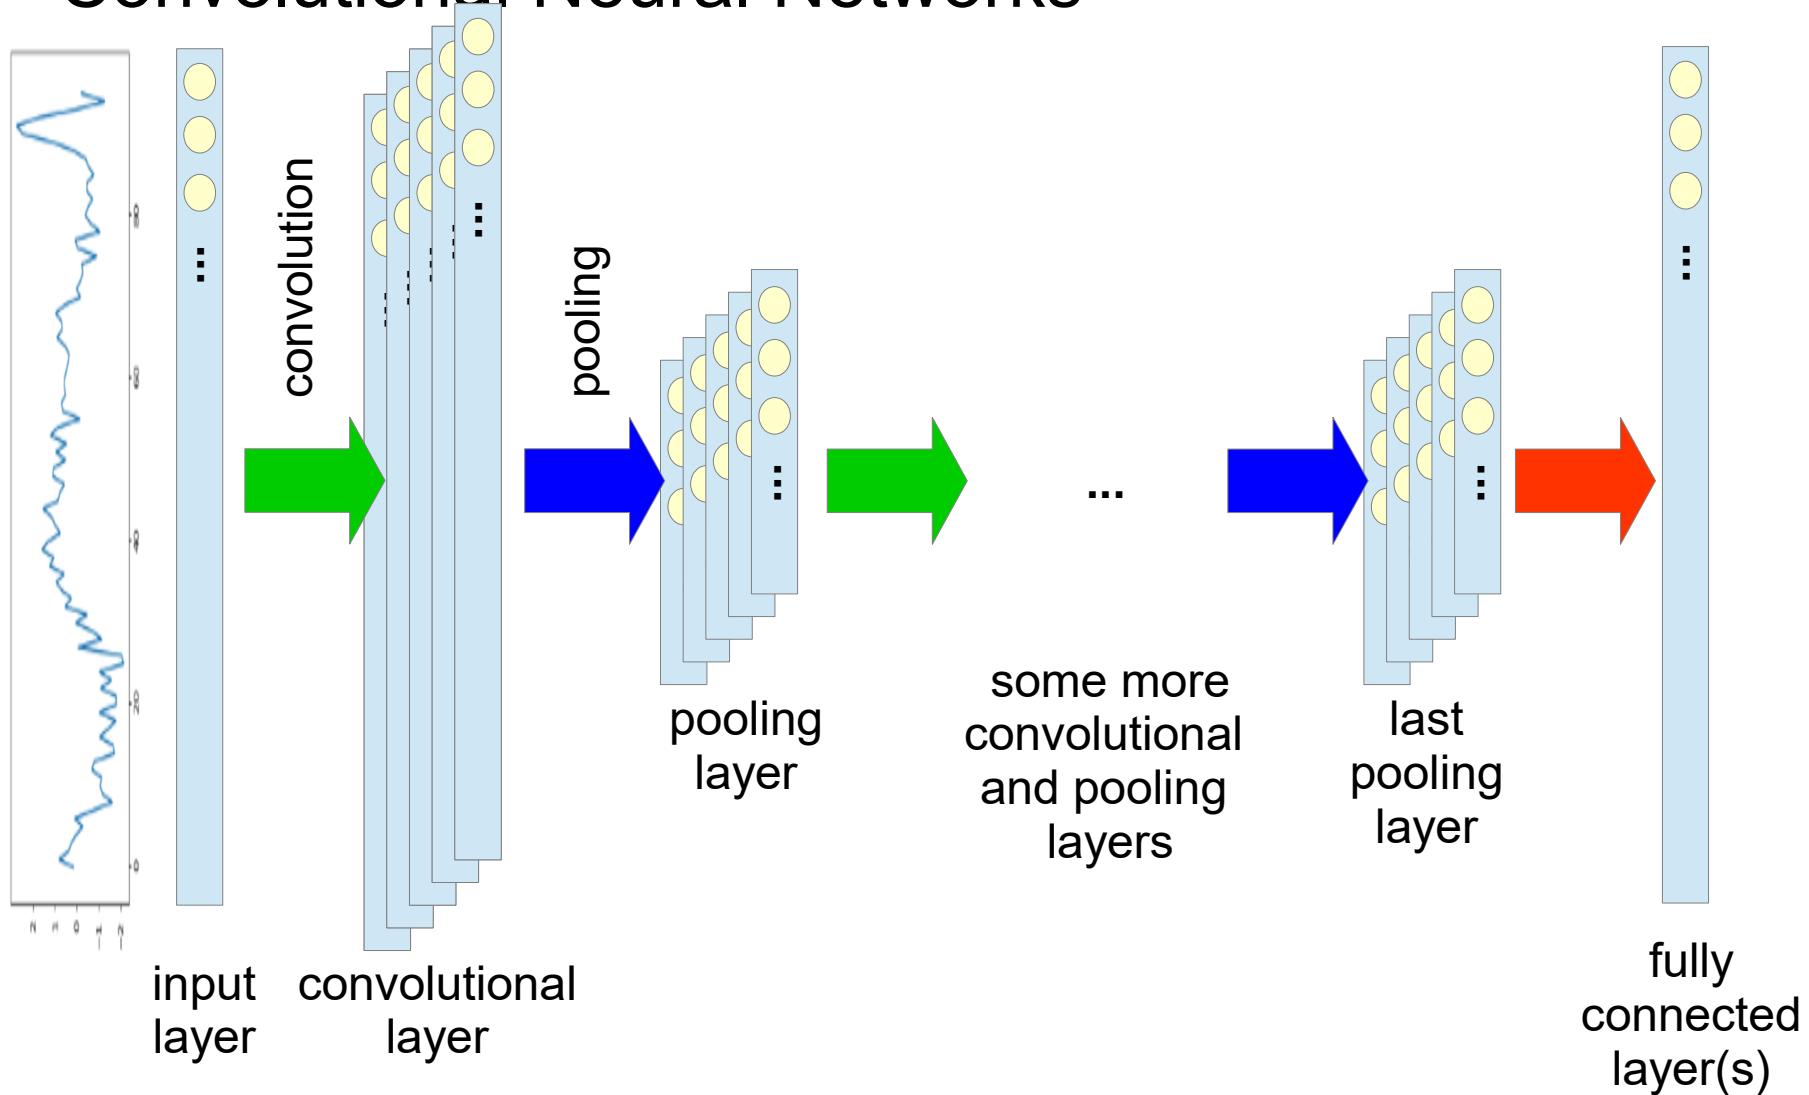
Convolutional Neural Networks



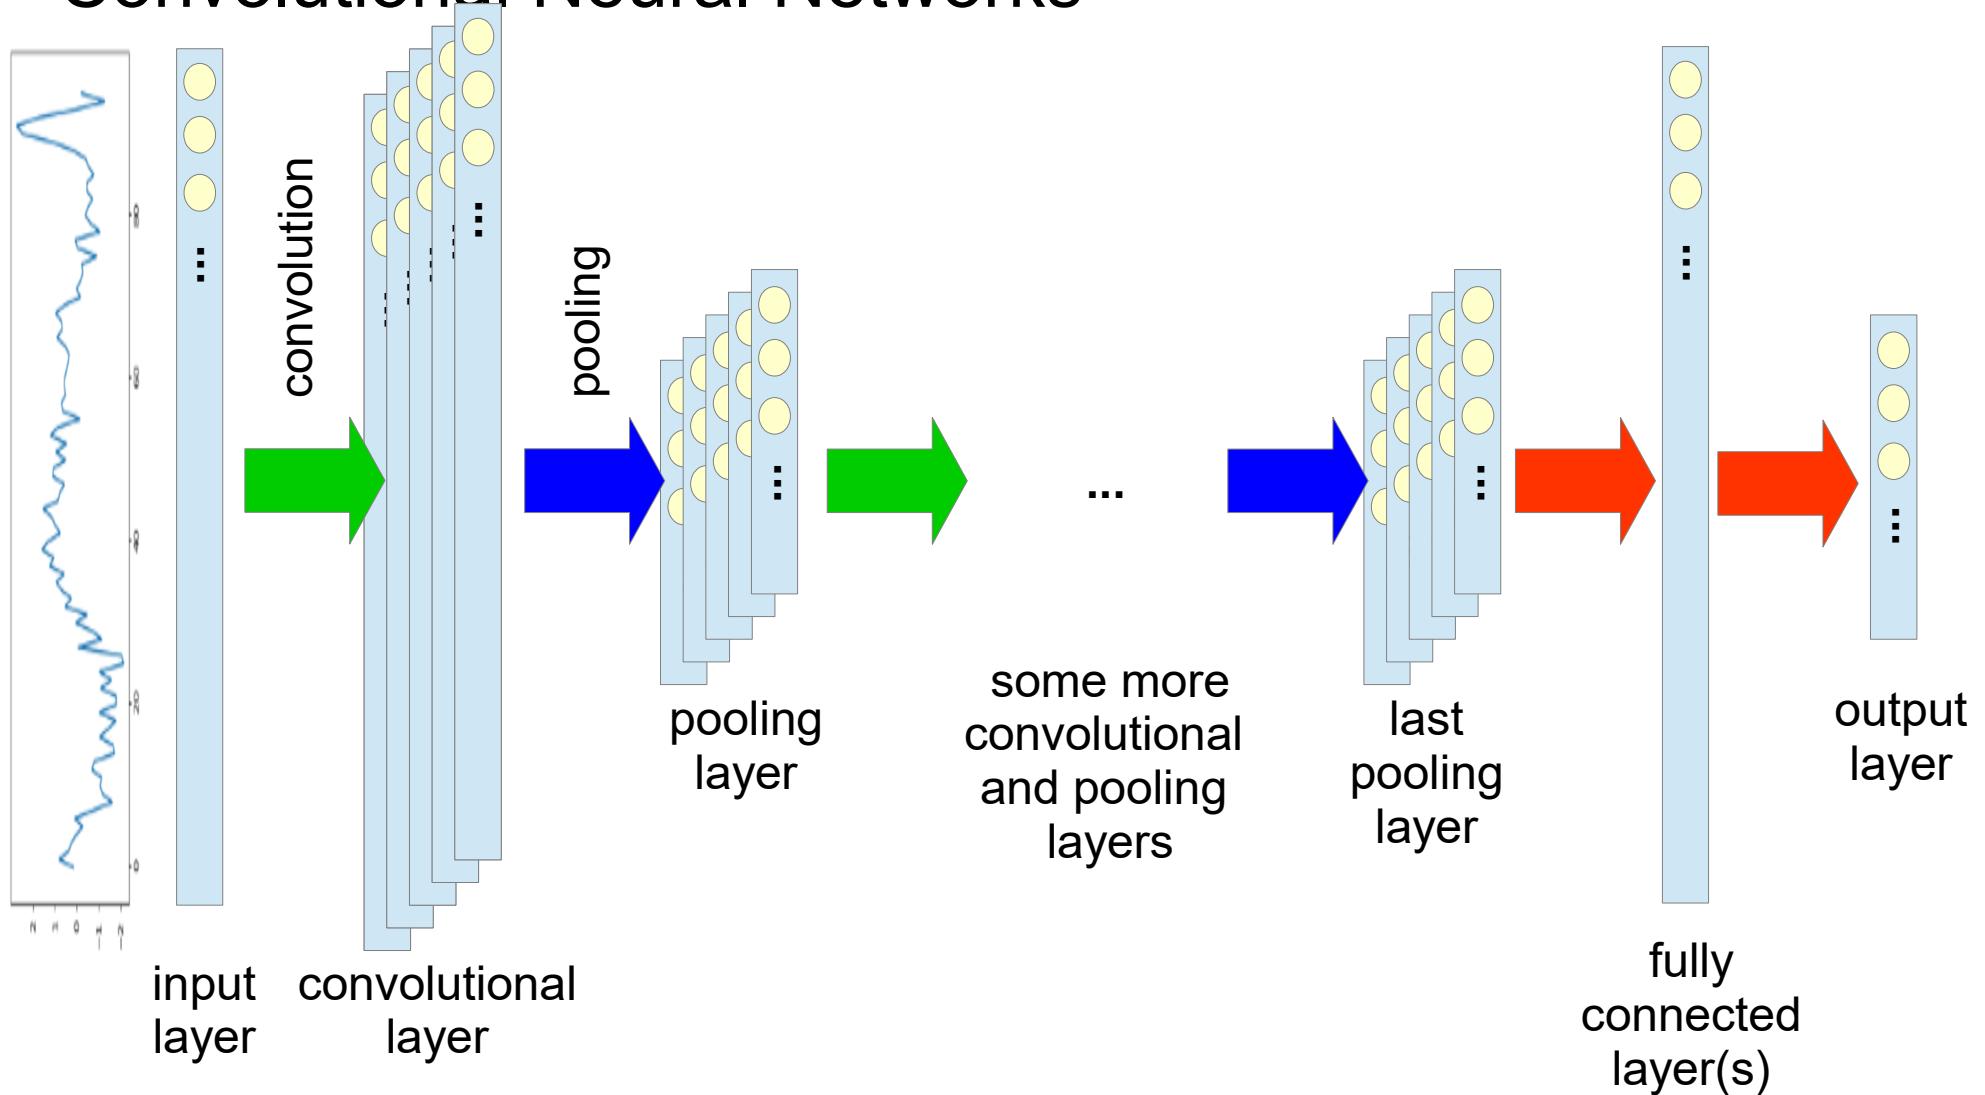
Convolutional Neural Networks



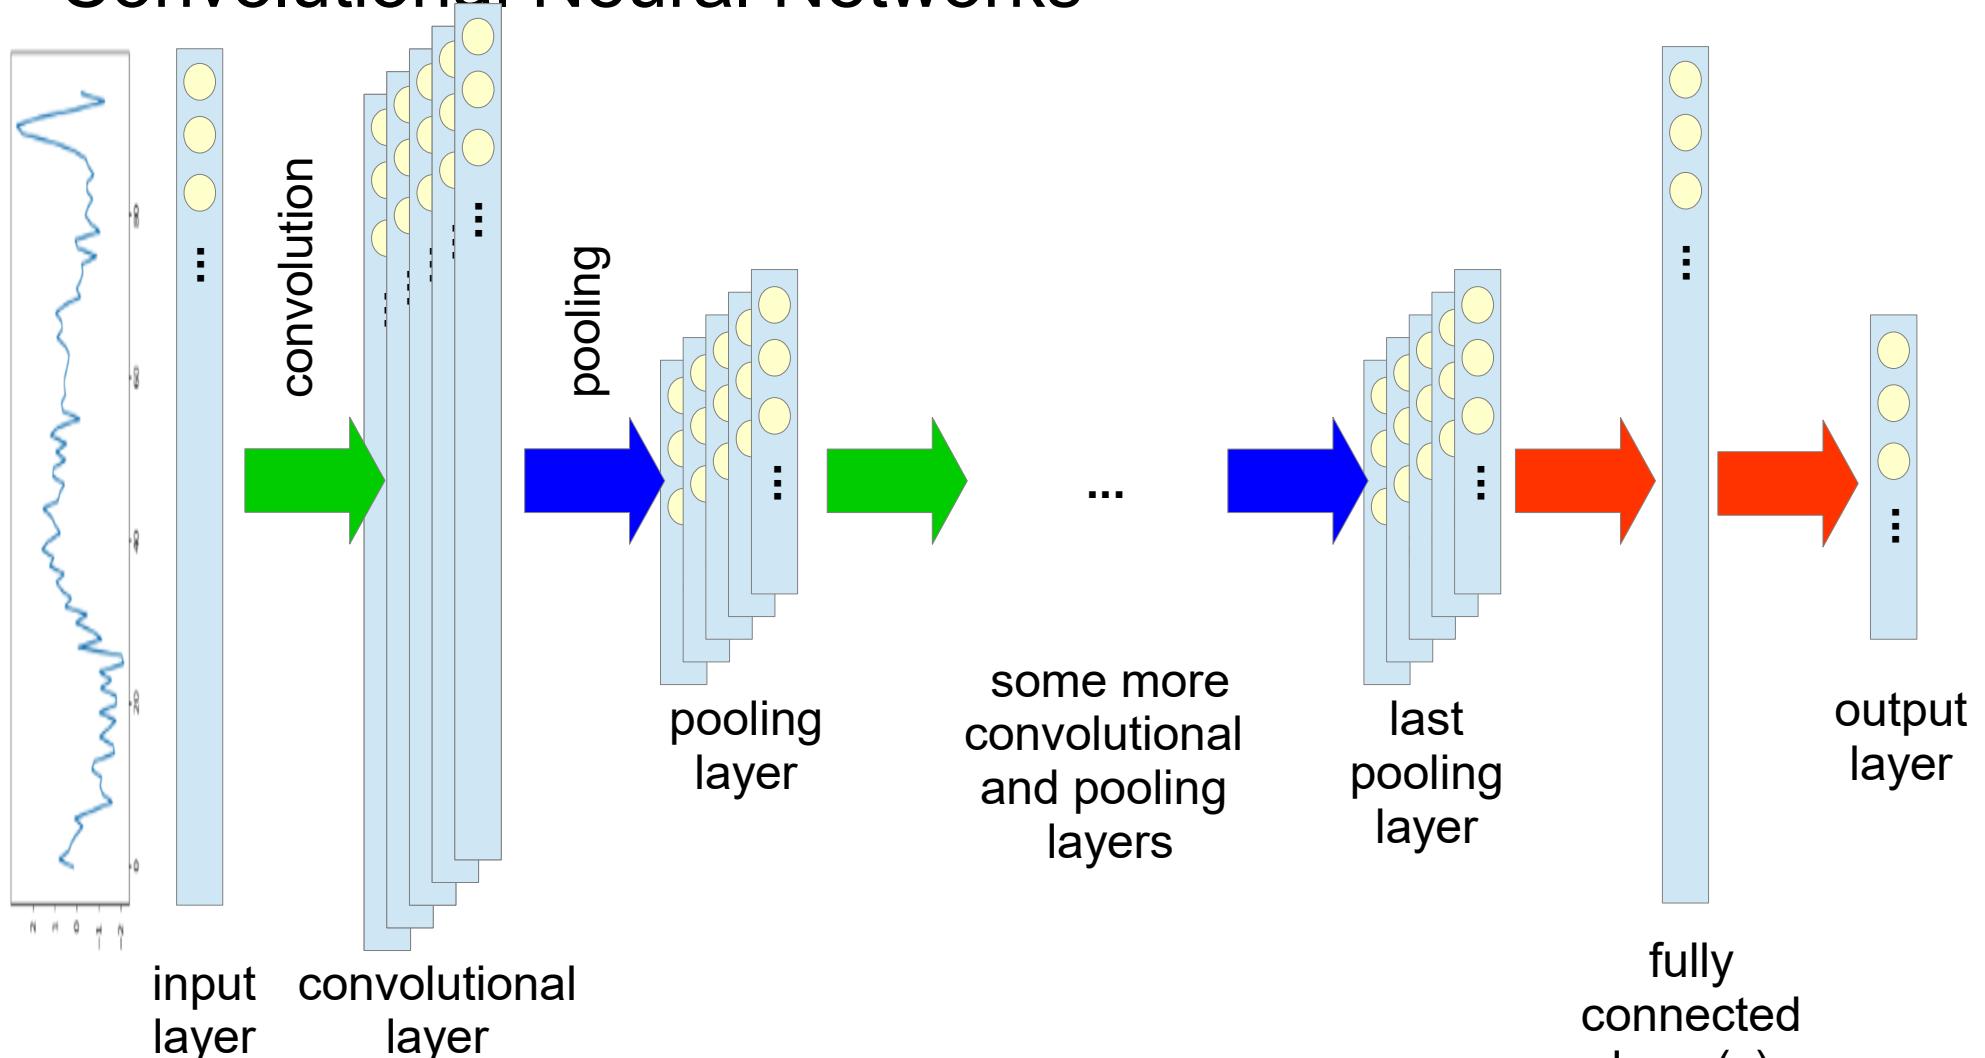
Convolutional Neural Networks



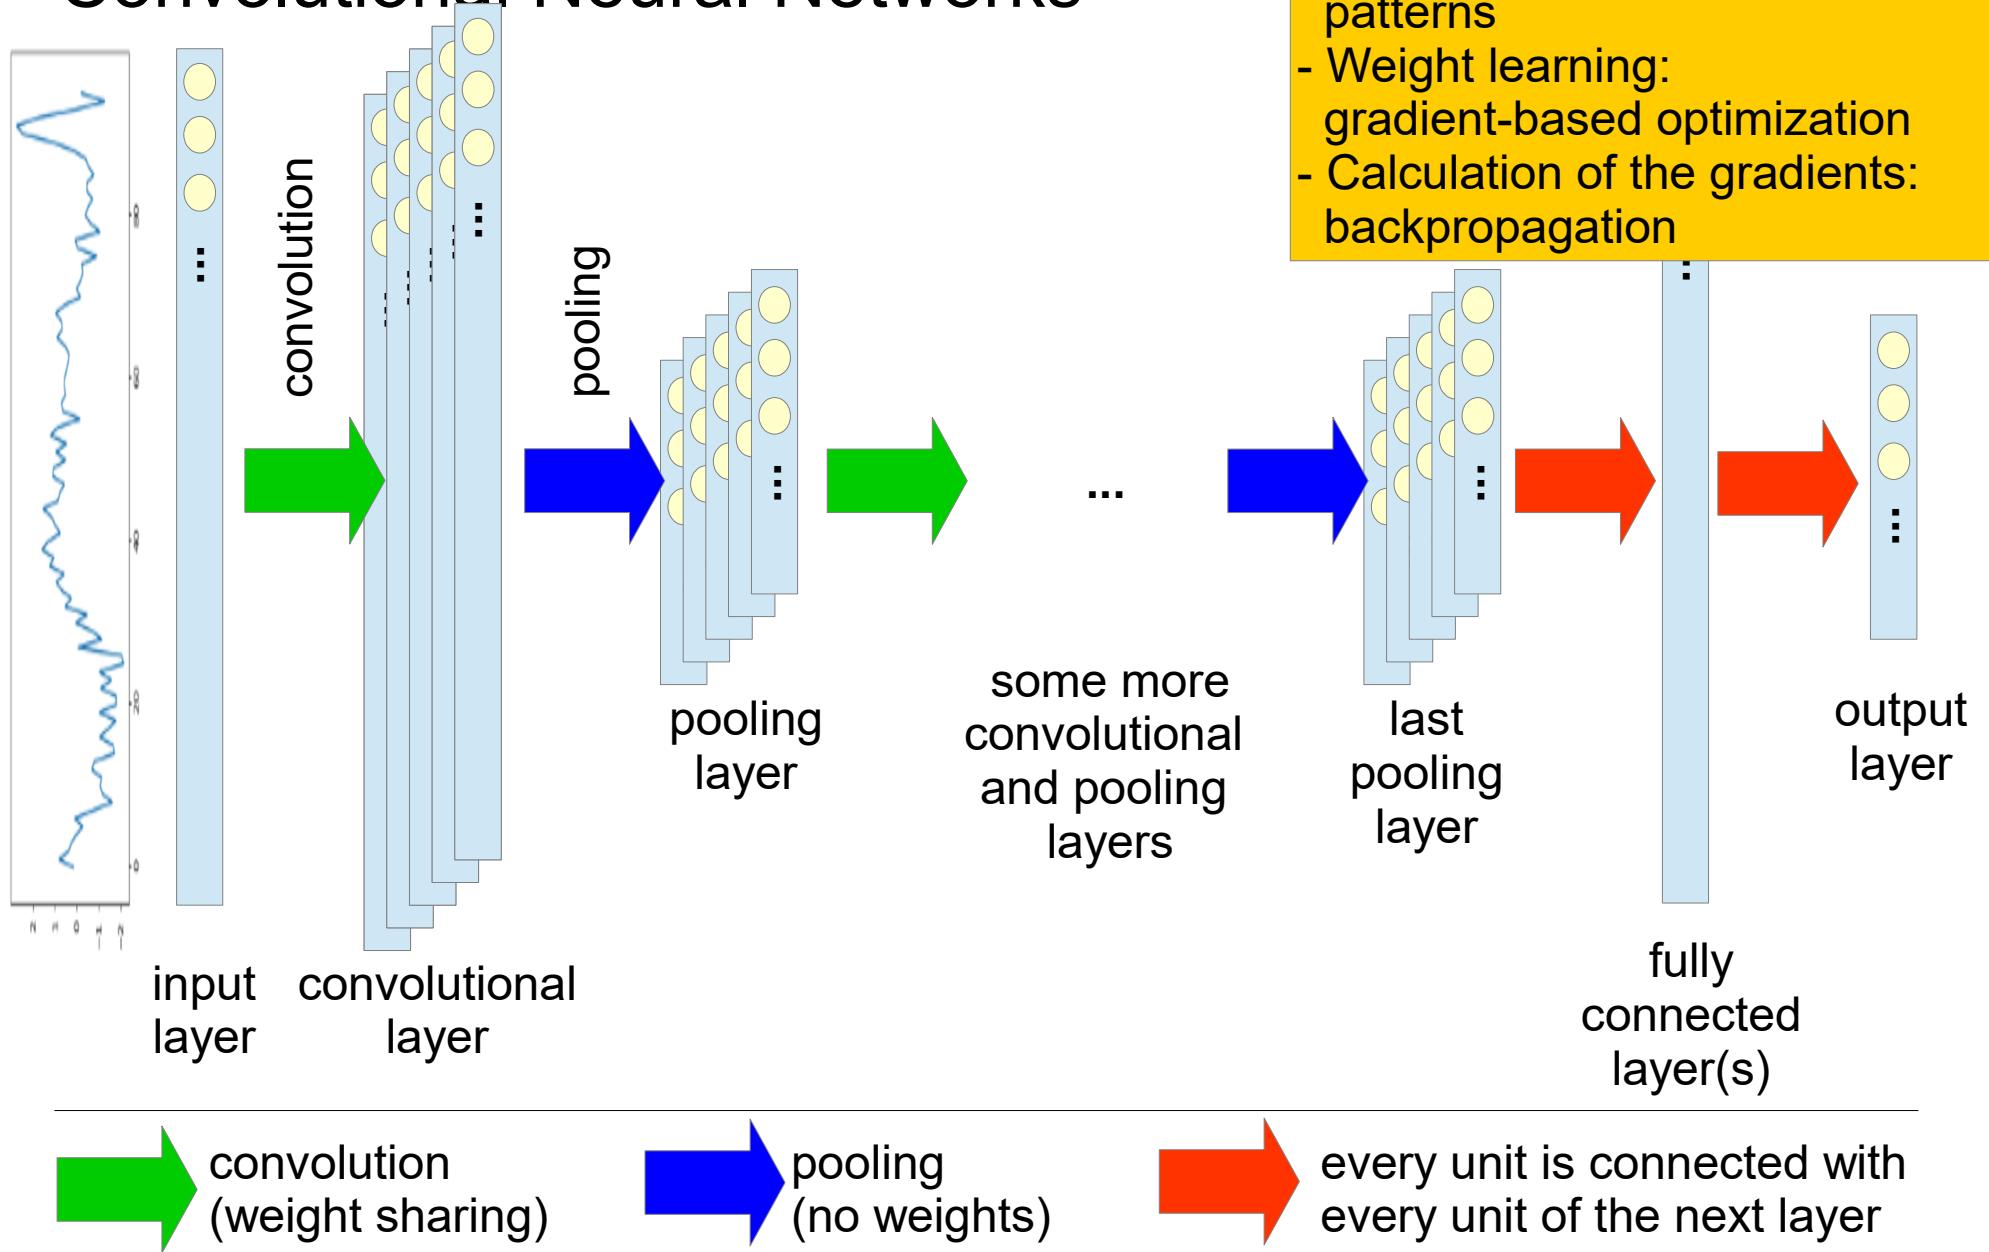
Convolutional Neural Networks



Convolutional Neural Networks



Convolutional Neural Networks



Classification based on Local Patterns

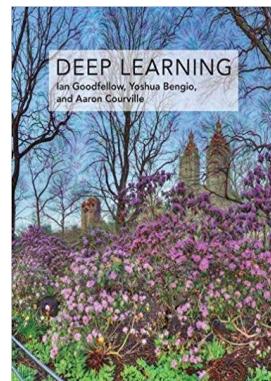
- Motif-based classification

Buza, Schmidt-Thieme (2009): Motif-based classification of time series with Bayesian networks and SVMs, Advances in Data Analysis, Data Handling and Business Intelligence. Springer, Berlin, Heidelberg, pp. 105-114

- Shapelet-based classification

Hills et al. (2014): Classification of time series by shapelet transformation, Data Mining and Knowledge Discovery, 28(4), pp. 851-881

- Convolutional Networks



Ian Goodfellow, Yoshua Bengio, Aaron Courville (2016):
Deep Learning, <http://www.deeplearningbook.org>

Implementation of Machine Learning Algorithms

Options

- 1. Implement everything from scratch
- 2. Implement training and prediction algorithm using basic functionality of software libraries such as dot product, matrix multiplications, etc.
- 3. Implement the model using more advanced functionality of software libraries, such as calculation of the gradients, optimization algorithms, built-in functions to train and evaluate the model
- 4. Integrate an existing model into your system

Some Implementation issues

1. Selection of the programming language, software libraries, development tools (IDE) and methods

Copyright, licence, unit tests, code review, coding and documentation standards, repository...

2. Availability of labeled data, data preprocessing

This may be a technological and legal issue simultaneously.

3. Optimization of computational costs, usage of RAM, energy consumption, availability of suitable hardware

Selection of data types, GPU, TPU, HPC, etc.

4. Other issues

Safety („cover your ass“),
compliance with law and ethical standards, etc.

Data Preprocessing (Data Cleaning)

- Handle Missing Values
- Elimination Duplicates
- Outlier Detection
- Handle Inconsistent Values
 - (e.g. replace misspelled values)

TensorFlow

- State-of-the-art machine learning library
- Suited for deep learning
- <https://www.tensorflow.org/>

Programming Homeworks

Programming Homeworks

- 1-3 students are expected to work in a team
(can be the same team as for the presentations, but it can be different as well)
- Three topics: Data Preprocessing (DP), Regression (R), Classification with Logistic Regression (C)
each team is expected to solve ALL tasks from ALL these topics
- Deadline: 4st Nov, **24 pts**, submission via e-mail to buza@inf.elte.hu
- You should deliver: code + contribution of team members

like this:

Zoltan Horvath	25%	DP-1, DP-2, R-1...
Krisztian Buza	30%	R-1, R-2, C-2...
Zakarya Farou	45%	R-1, R-3, R-4, R-5, C-1...

- In the lecture on the 6th Nov, for each task T, a randomly chosen team will be selected. If the selected team delivered a solution for T, but can not explain their solution, team members' scores (for the entire programming homework) will be set to 0 (zero).

Programming Homeworks

- You will need
 - ... a Python interpreter / IDE, such as
 - a command line interpreter
 - Anaconda Jupyter Notebook
 - Google Collab
 - ... to install some libraries (numpy, scikit-learn...)
- Tasks are provided in two formats: html and ipynb
 - Feel free to use which is more convinient for you
- If you use the ipynb files, you may need to rename the **ipynb** files from *.html to *.ipynb (if your browser saves them as html files)
- **Please google for installation/usage instructions**

Programming Homework

- Most of the code is already prepared, you only need to complete some of the functions and execute the completed code.
- **Be honest:** the solution of similar tasks are available on the internet. Studying these solutions is allowed, because the goal of the homework is not that you suffer, but that you learn.
- **However:** you should **understand** the solutions that is submitted, and you should not ask anyone else (who is not a member of your team) to solve the tasks instead of you.

Summary

Summary

- Overview of the techniques known under the umbrella of „deep learning“
- New mathematical operation: Convolution
- Convolutional Networks
- Issues related to the implementation of machine learning algorithms

Essential Concepts

- GPU, TPU („AI chip“)
- Pre-training
 - Supervised pre-training
 - Unsupervised pre-training
- Drop-out
- Convolution
- Convolutional networks
- Data Augmentation
- Data Preprocessing, Data Cleaning